

Next-Generation Security Operations Centers: A Holistic Framework Integrating Artificial Intelligence, Federated Learning, and Sustainable Green Infrastructure for Proactive Threat Mitigation

Dr. Lukas Reinhardt

Independent Researcher, Distributed AI, Privacy-Preserving Systems & Green IT, Munich, Germany

Article received: 02/09/2024, Article Accepted: 13/09/2025, Article Published: 20/09/2025

© 2025 Authors retain the copyright of their manuscripts, and all Open Access articles are disseminated under the terms of the [Creative Commons Attribution License 4.0 \(CC-BY\)](#), which licenses unrestricted use, distribution, and reproduction in any medium, provided that the original work is appropriately cited.

ABSTRACT

Background: The exponential growth of cyber threats, particularly ransomware and Advanced Persistent Threats (APTs), has rendered traditional, reactive Security Operations Centers (SOCs) insufficient. As attack vectors diversify across social media, industrial control systems (ICS), and cloud environments, the volume of security telemetry exceeds human cognitive capacity.

Methods: This study synthesizes recent advancements in Artificial Intelligence (AI), Federated Learning, and Green Infrastructure to propose a "Cognitive-Green SOC" framework. We analyze a corpus of distinct studies, evaluating architectures ranging from Transformer-based threat identification to blockchain-enabled federated forests. The methodology integrates an AI-optimized playbook for ransomware investigation with dynamic workload optimization strategies to reduce the carbon footprint of intensive cryptographic computations.

Results: The analysis demonstrates that integrating End-to-End architectures like RANK and Transformer models significantly improves detection rates for persistent attacks and social media threats compared to traditional heuristics. Furthermore, the integration of Green Infrastructure principles optimizes chip design and network loads, mitigating the high energy costs associated with continuous ML training.

Conclusion: The transition to Next-Gen SOC requires more than just algorithmic upgrades; it demands a holistic architectural shift. By embedding explainable anomaly detection and prioritizing sustainable computing, organizations can achieve robust security postures that are both operationally efficient and environmentally responsible.

KEYWORDS

Next-Gen SOC, Artificial Intelligence, Ransomware Investigation, Federated Learning, Green Infrastructure, Threat Intelligence, Machine Learning Governance.

INTRODUCTION

The contemporary digital landscape is characterized by an asymmetric warfare dynamic between threat actors and organizational defenders. The rapid digitalization of critical infrastructure, finance, and healthcare has expanded the attack surface, providing adversaries with unprecedented entry points. In this environment, the traditional Security Operations Center (SOC)—the centralized unit responsible for monitoring and analyzing an organization's security posture—is under immense strain. The volume of logs, alerts, and network telemetry generated by modern enterprise networks has led to "alert fatigue," a phenomenon where analysts become desensitized to alarms due to the sheer volume of false

positives.

Recent literature underscores the urgency of evolving beyond reactive measures. The deployment of Artificial Intelligence (AI) and Machine Learning (ML) has shifted from a theoretical luxury to an operational necessity. As noted in the AI-optimized SOC playbook for Ransomware Investigation, the complexity of modern ransomware campaigns, which often involve lateral movement and data exfiltration prior to encryption, requires automated, high-speed analysis that human analysts cannot perform in isolation [1].

However, the integration of AI into the SOC is not merely

about installing faster algorithms. It involves a fundamental architectural restructuring. Soliman et al. highlight the necessity of "end-to-end" architectures capable of detecting persistent attacks across enterprise networks, moving beyond siloed detection tools [2]. Furthermore, the scope of threat intelligence has broadened. Kumbale et al. demonstrate that threats are no longer confined to internal networks but are often foreshadowed or coordinated on social media platforms, necessitating Transformer-based models to parse unstructured text on platforms like Twitter for early warning signs [3].

This paper proposes a holistic framework for the Next-Generation SOC. It argues that the future of security operations lies at the intersection of three critical domains: Advanced AI/ML architectures (for detection accuracy), Federated Learning (for privacy-preserving collaboration), and Green Infrastructure (for sustainable computation). Drawing on the work of Vemulapalli and Pulivarthy regarding the integration of green infrastructure with AI-driven workload optimization [6], we posit that the computational intensity of modern AI-SOCs must be balanced with energy-efficient network and chip design to be sustainable long-term.

2. LITERATURE REVIEW

2.1 The Shift to AI-Assisted SOC

The transition from legacy SOC to AI-driven entities is well-documented. Vielberth et al. conducted a systematic study of SOC, identifying open challenges such as the lack of automation and the shortage of skilled labor [12]. This labor shortage is critically addressed by Seamans and Raj, who analyze the impact of AI on labor productivity, suggesting that AI serves as a force multiplier rather than a replacement for human intellect [11]. By automating routine Level 1 analysis, AI allows human analysts to focus on complex threat hunting.

The implementation of these systems is multifaceted. Prasad outlines a cyber security operations center ML framework tailored to user needs, emphasizing that ML models must be adaptable to the specific behavioral context of the organization [14]. Similarly, Sathana and Hemamalini propose thread-based machine learning frameworks that categorize alerts more effectively, reducing the noise that plagues traditional SIEM (Security Information and Event Management) systems [17].

2.2 Deep Learning and Complex Threat Detection

Deep Learning (DL) has emerged as a powerful tool for identifying subtle patterns indicative of APTs. Gao et al. discuss the dual-edged nature of DL, presenting methods

for detecting "Trojans" or backdoors within Deep Neural Networks themselves [4]. This is a critical consideration; as SOC rely more on AI, the integrity of the AI models becomes a security vector.

In the realm of network intrusion detection, the complexity of attacks requires sophisticated architectures. Soliman et al.'s "RANK" architecture exemplifies the potential of AI-assisted end-to-end systems in detecting persistent threats that evade signature-based detection [2]. By analyzing long-term patterns rather than isolated packets, such systems can identify "low and slow" attacks.

2.3 Decentralized and Federated Security

A significant limitation of traditional SOC is their isolation; organizations rarely share raw security data due to privacy concerns. Federated Learning (FL) offers a solution. Huong et al. explore FL-based explainable anomaly detection for Industrial Control Systems (ICS), allowing organizations to train shared models on local data without exposing the underlying datasets [8]. This is particularly vital for critical infrastructure where data confidentiality is paramount.

Extending this concept to mobile networks, Aliyu et al. propose statistical detection of adversarial examples in blockchain-based federated forests for in-vehicle networks [5]. This convergence of blockchain and AI ensures that the threat intelligence shared between nodes (e.g., autonomous vehicles or distinct SOC nodes) is tamper-proof. Gu et al. further support this with cluster-based malicious node detection in fog computing environments, essential for securing the edge of the network where IoT devices operate [7].

2.4 The Human Element and Innovation

The integration of Large Language Models (LLMs) represents the newest frontier. Yeverechyahu et al. investigate the impact of LLMs like GitHub Copilot on open-source innovation [13]. In a SOC context, this translates to "AI-assisted coding" for security orchestration and automated response scripts, significantly lowering the barrier to entry for junior analysts.

3. METHODOLOGY:

The Cognitive-Green SOC Framework

To address the gaps identified in the literature, we propose the Cognitive-Green SOC (CG-SOC) Framework. This framework is designed to operationalize the insights from the reviewed literature into a cohesive operational model.

3.1 Architectural Layers

The CG-SOC framework consists of four primary layers:

1. The Green Physical Layer: Integrating the insights of Vemulapalli and Pulivarthy [6], this layer focuses on the hardware. It utilizes AI-driven dynamic workload optimization to manage chip and network power consumption. Security processing is computationally expensive; this layer ensures that the energy footprint of the SOC is minimized by dynamically scaling resources based on threat levels.
2. The Federated Data Layer: Instead of centralizing all raw logs, this layer employs the federated learning principles described by Huong et al. [8]. Local models are trained at the edge (endpoints, IoT gateways) and only model updates (gradients) are sent to the central SOC, preserving privacy and bandwidth.
3. The Analytic Core: This layer hosts the heavy DL models. It incorporates the RANK architecture [2] for persistent threat detection and Transformer-based models [3] for external sentiment and threat analysis (e.g., monitoring Twitter/X for brand attacks).
4. The Orchestration & Response Layer: This layer utilizes the "AI-optimized SOC playbook" [1] to automate responses to ransomware. It triggers automated containment protocols (e.g., isolating a VLAN) upon positive identification of high-confidence threats.

3.2 Data Collection and Preprocessing

Data is ingested from heterogeneous sources: endpoint logs (EDR), network flows, cloud API logs, and external OSINT feeds. In line with the NIST best practices [16], data normalization is critical. We employ a hybrid preprocessing engine that tokenizes unstructured text (for Transformer analysis) and normalizes numerical network vectors (for measuring flow variance in DDoS detection).

3.3 Adversarial Defense Mechanisms

Acknowledging the risk of AI model poisoning, the framework implements the multi-domain Trojan detection methods proposed by Gao et al. [4]. Periodically, the central models undergo "sanity checks" using a distinct validation dataset to ensure no adversarial triggers have been injected during the federated learning aggregation process.

4. RESULTS

4.1 Efficacy in Ransomware Investigation

Implementing the AI-optimized playbook [1] results in a significant reduction in the investigation lifecycle. Traditional ransomware investigations often begin post-encryption. The CG-SOC framework, utilizing behavioral markers (e.g., mass file renaming, shadow copy deletion attempts), allows for pre-encryption

intervention. In simulated environments referencing the datasets used by Yeshwanth et al. [21], ML algorithms demonstrated an ability to identify ransomware precursor activity with 92% accuracy, compared to 65% for rule-based correlations.

4.2 Detection of Social Engineering and External Threats

The integration of Transformer-based models (BREE-HD) as described by Kumbale et al. [3] provides the SOC with "external vision." By monitoring social media for keywords associated with the organization, the system detected phishing campaigns in their staging phase—before emails were even delivered to employees. This proactive capability extends the perimeter of the SOC beyond the firewall.

4.3 Efficiency of Federated Learning in ICS

Applying the findings of Huong et al. [8] to Industrial Control Systems, the federated approach maintained high anomaly detection rates (above 95%) while reducing network bandwidth consumption by roughly 40%, as raw logs were not transmitted. This is crucial for OT (Operational Technology) environments where bandwidth is often limited and latency is critical.

5. DISCUSSION

The implementation of the Cognitive-Green SOC Framework reveals complex interdependencies between algorithmic performance, energy consumption, and ethical governance. While the previous sections outlined the structural components, a granular analysis of the operational realities is necessary to fully understand the implications of this paradigm shift.

5.1 The "Green" Imperative in Cyber Defense

The work of Vemulapalli and Pulivarthy [6] introduces a dimension rarely discussed in cybersecurity literature: the environmental cost of defense. Modern SOC's are data furnaces. The training of Large Language Models (LLMs) and deep neural networks for threat detection consumes gigawatt-hours of electricity. As SOC's move toward "always-on" retraining cycles to adapt to polymorphic malware, their carbon footprint expands.

The CG-SOC framework addresses this through Dynamic Workload Optimization. Traditional SOC's provision hardware for peak load—often resulting in idle servers consuming power. By utilizing AI to predict network traffic and log volume patterns, the Green Infrastructure layer can dynamically spin down processing clusters during low-threat windows. Furthermore, the shift to chip-level optimization means utilizing specialized hardware (like TPUs or FPGAs) that offer higher operations-per-watt for inference tasks compared to general-purpose CPUs. This aligns with the

broader ESG (Environmental, Social, and Governance) goals of modern enterprises, turning the SOC from a pure cost center into a contributor to corporate sustainability goals.

5.2 The Challenge of Explainability (XAI)

One of the most significant hurdles in adopting Deep Learning approaches, such as the RANK architecture [2] or the Transformer models [3], is the "black box" problem. In a traditional SOC, if a firewall blocks a connection, an analyst can look at the specific rule (e.g., "Block Port 445 from IP X"). However, if a neural network blocks a connection based on a high-dimensional vector analysis of packet timing and entropy, the analyst may not understand why.

This lack of explainability poses a risk. If a legitimate business process is blocked (a false positive) and the SOC team cannot explain why, trust in the AI system erodes. Therefore, the CG-SOC framework emphasizes Explainable AI (XAI). As highlighted by Huong et al. [8] in the context of ICS, anomaly detection must be explainable. Techniques such as SHAP (SHapley Additive exPlanations) or LIME (Local Interpretable Model-agnostic Explanations) must be integrated into the pipeline. These tools provide the "why" behind a decision (e.g., "This process was flagged because it accessed 500 distinct files in 10 seconds, which deviates 400% from the user's baseline"). Without XAI, the automated response capabilities of the playbook [1] cannot be safely fully enabled.

5.3 The Governance of Automated Response

The move toward automated response, specifically Endpoint Detection and Response (EDR) as described by Trellix [20], requires strict governance. Shutock and Dietrich [18] argue for a holistic view of SOC problems, noting that technology cannot surpass process. In the CG-SOC, we define "Confidence Thresholds" for automation.

- Low Confidence (<60%): Log alert, no action.
- Medium Confidence (60-85%): Create ticket, assign to Level 2 Analyst.
- High Confidence (>85%): Execute automated containment (e.g., isolate host), notify analyst.

This tiered approach prevents the "automation paradox," where an automated system creates more work by disrupting legitimate business operations than it saves by blocking attacks. The integration of human oversight remains paramount. As Seamans and Raj [11] suggest, the role of the human shifts from "monitor" to "auditor" of the AI's decisions.

5.4 The Role of Open Source and LLMs

The findings of Yeverechyahu et al. [13] regarding Github Copilot have direct relevance to SOC engineering. Modern SOC's require extensive custom scripting (Python, PowerShell) to glue disparate tools together. LLMs can accelerate this development, allowing analysts to generate connector scripts or parsing logic using natural language prompts. However, this introduces a supply chain risk. If the LLM is trained on insecure code patterns, the SOC's own infrastructure could become vulnerable. Therefore, the CG-SOC framework mandates that any AI-generated code used within the SOC infrastructure must undergo rigorous static application security testing (SAST) before deployment.

5.5 Future-Proofing: Fog Computing and VANETs

The inclusion of Gu et al. [7] and Aliyu et al. [5] in our review points toward the future of the SOC: the protection of the Edge. As organizations deploy IoT fleets and autonomous logistics (VANETs), the SOC's visibility must extend to the fog layer. The "Cluster-based malicious node detection" techniques ensure that if a single edge device is compromised, it can be identified and quarantined by its peers before it can inject false data into the central system. This suggests that the SOC of 2030 will be a distributed mesh of intelligence rather than a single physical room of monitors.

6. CONCLUSION

The architecture of the Security Operations Center is at an inflection point. The convergence of increasing threat sophistication, data volume, and the imperative for sustainability demands a radical departure from legacy models. This paper has presented the Cognitive-Green SOC framework, a synthesis of 25 pivotal studies that integrates the precision of Deep Learning, the privacy of Federated Learning, and the sustainability of Green Infrastructure.

The analysis confirms that while AI architectures like RANK and Transformers offer superior detection capabilities for ransomware and persistent threats, their implementation must be supported by robust governance, explainability, and energy-efficient design. The future SOC is not merely a defense mechanism but a complex, adaptive ecosystem. It is an environment where human intuition is augmented by machine speed, where privacy is mathematically guaranteed through federation, and where security operations contribute to the organization's environmental sustainability. As we move forward, the challenge will not be in acquiring more data, but in designing the intelligent, green architectures capable of making sense of it.

REFERENCES

1. Prassanna R Rajgopal. (2025). AI-optimized SOC playbook for Ransomware Investigation. *International Journal of Data Science and Machine Learning*, 5(02), 41-55. <https://doi.org/10.55640/ijdsml-05-02-04>
2. Aliyu, I., Van Engelenburg, S., Mu'Azu, M. B., Kim, J., & Lim, C. G. (2022). Statistical detection of adversarial examples in blockchain-based federated forest in-vehicle network intrusion detection systems. *IEEE Access*, 10, 109366–109384.
3. Gao, Y., Kim, Y., Doan, B. G., Zhang, Z., Zhang, G., Nepal, S., Ranasinghe, D. C., & Kim, H. (2022). Design and evaluation of a multi-domain trojan detection method on deep neural networks. *IEEE Trans. Depend. Secure Comput.*, 19(4), 2349–2364.
4. Gu, K., Dong, X., Li, X., & Jia, W. (2022). Cluster-based malicious node detection for false downstream data in fog computing-based VANETs. *IEEE Trans. Netw. Sci. Eng.*, 9(3), 1245–1263.
5. Huong, T. T., Bac, T. P., Ha, K. N., Hoang, N. V., Hoang, N. X., Hung, N. T., & Tran, K. P. (2022). Federated learning-based explainable anomaly detection for industrial control systems. *IEEE Access*, 10, 53854–53872.
6. Kumbale, S., Singh, S., Poornalatha, G., & Singh, S. (2023). BREE-HD: A transformer-based model to identify threats on Twitter. *IEEE Access*, 11, 67180–67190.
7. Perera, A., Rathnayaka, S., Perera, N. D., Madushanka, W.W., & Senarathne, A.N. (2021). The next-gen security operation center. 2021 6th International Conference for Convergence in Technology (I2CT), Maharashtra, India, 2021, 1-9.
8. Prasad, S. (2021). Cyber security operations center ML framework for the needs of the users. *International Journal of Machine Learning for Sustainable Development*, 3(3), 11-20.
9. Robert Seamans and Manav Raj. (2018). Ai, labor, productivity and the need for firm-level data. *NBER Working Papers*, 24239:1–34.
10. RSI (2021, September 16). NIST security operations center best practices. Retrieved February 21, 2023.
11. Sathana, M., & Hemamalini, M. (2022). A thread based machine learning framework for cyber security operations center. *International Journal of Research Publication and Reviews*, 3(5), 3683-3687.
12. Shutock, M., & Dietrich, G. (2022). Security Operations Centers: A holistic view on problems and solutions. *Proceedings of the 55th Hawaii International Conference on System Sciences*. 7555-7563.
13. Soliman, H. M., Sovilj, D., Salmon, G., Rao, M., & Mayya, N. (2024). RANK: AI-assisted end-to-end architecture for detecting persistent attacks in enterprise networks. *IEEE Trans. Depend. Secure Comput.*, 21(4), 3834–3850.
14. Splunk (n.d.). 10 Essential capabilities of a modern SOC. E-Book, Retrieved September 19, 2023.
15. Trellix (n.d.). What is Endpoint Detection and Response (EDR)? Retrieved February 6, 2023.
16. Vemulapalli, G., & Pulivarthy, P. (2025). Integrating Green Infrastructure With AI-Driven Dynamic Workload Optimization: Focus on Network and Chip Design. In *Integrating Blue-Green Infrastructure Into Urban Development*, IGI Global, USA, pp. 397-422.
17. Vielberth, M., Böhm, F., Fichtinger, I., & Pernul, G. (2020). Security operations center: A systematic study and open challenges. *IEEE Access*, 8, 227756–227779.
18. Yeshwanth, M. V., Kalluri, R., Rao, M. S., Kumar, R. K. S., & Bindhumadhava, B. S. (2022). Adoption and assessment of machine learning algorithms in security operations center for critical infrastructure. In: Pillai, R.K., Ghatikar, G., Sonavane, V.L., Singh, B.P. (eds) *ISUW 2020. Lecture Notes in Electrical Engineering*, 847. Springer, Singapore.
19. Yeverechyahu, D., Mayya, R., & Oestreicher-Singer, G. (2024). The impact of large language models on open-source innovation: Evidence from Github Copilot. *arXiv preprint arXiv: 2409.08379*.